
HOMework

Solutions

1. (33 points) **Tree Cutting.** You are given a tree of n vertices, rooted at vertex 0. Consider the following game: every turn, you pick a vertex uniformly at random from the remaining tree and cut the node, together with the node's subtree. The procedure stops when the whole tree is cut (i.e., the root is cut).

Design an algorithm that computes the expected total number of cuts when the procedure stops, in linear time.

Note. If you are able to solve the case where the graph is a line: $0 \rightarrow 1 \rightarrow 2 \rightarrow \dots \rightarrow n-1$ in at most linear time, you will get 20 points. This can be a good starting point.

Hint. (Linearity of expectation). Define e_v as a random variable indicating whether v is ever chosen before the whole tree is cut.

First, let us look at the simpler case where the graph is a line: $0 \rightarrow 1 \rightarrow 2 \rightarrow \dots \rightarrow n-1$.

Like the course material, we define e_v , $e_v = 1$ if node v is selected before any of its ancestors (nodes 0 to $v-1$), and $e_v = 0$ otherwise. The objective, which is the expected total number of cuts required to remove the entire tree, is lists as follows:

$$E = \mathbb{E} \left[\sum_{v=0}^{n-1} e_v \right].$$

The reason that we can calculate the objective in such way is because for each node, we either make it responsible for one of these cuts or it is already been removed as part of a subtree when an ancestor is cut.

Now we need to calculate $\mathbb{E}[e_v]$ for each node:

$$E = \mathbb{E} \left[\sum_{v=0}^{n-1} e_v \right] = \sum_v \mathbb{E}[e_v].$$

Under the simpler case, we have the circumstance that the set of node from root 0 to node v , inclusive, has $\text{depth}(v)+1$ nodes. Since each node is equally likely to be selected at each turn (uniform random selection), the probability that node v is the first among these $v+1$ nodes to be selected is:

$$\mathbb{P}(e_v = 1) = \frac{1}{v+1},$$

which means, the probability that node v is removed by ancestors is:

$$\mathbb{P}(e_v = 0) = 1 - \frac{1}{v+1}.$$

As a result, the expected total number of cuts is:

$$E = \sum_{v=0}^{n-1} \mathbb{E}[e_v] = \sum_{v=0}^{n-1} \frac{1}{v+1}.$$

It is easily seen that the time we need to calculate E is $O(N)$.

Second, let us look at the general case. We now have an arbitrary rooted tree with n nodes. For each node v , we define $\text{depth}(v)$ as its depth (the number of edges from the root 0 to node v).

Same as the simpler case, we define e_v , $e_v = 1$ if node v is selected before any of its ancestors (nodes 0 to $v-1$), and $e_v = 0$ otherwise.

The objective is also the same:

$$E = \mathbb{E} \left[\sum_{v=0}^{n-1} e_v \right].$$

To calculate the probability that a node would be selected to remove trees, which means not to be removed because of its ancestors, now we only need to consider the path from node 0 to node v , in other words, this node to be selected under this path:

$$\mathbb{P}(e_v = 1) = \frac{1}{\text{depth}(v) + 1}.$$

To explain more specifically, since these $\text{depth}(v)+1$ nodes(ancestors and v) are the only ones that can prevent v from being selected first in its ancestral line, and selection is uniform. *Note: That is exactly what I mean by path :)*

As a result, the expected total number of cuts is:

$$E = \sum_v \mathbb{E}[e_v] = \sum_v \frac{1}{\text{depth}(v) + 1}.$$

Now the algorithm is also obvious, we just need to pre-compute the depth of each nodes and preform the *Linearity of expectation*. I list the steps as follows:

1. Perform DFS to computes depth for all nodes.
2. For each node, compute $\frac{1}{\text{depth}(v)+1}$ and sum them up.
3. The sum obtained is the expected number of cuts required to remove the entire tree.

All the steps is $O(N)$ so we get a linear time approach.

-
2. (33 points) **Picking Teams.** You have n people that you want to divide into two teams, A and B . Every person needs to be included in exactly one of the teams. The i -th person has two associated prices:

- a_i – the price to include them in team A
- b_i – the price to include them in team B

You also learn that there are m pairs of friends among these n people, which are specified by (u_j, v_j, c_j) ¹ where we guarantee that $u_j \neq v_j$. If persons u_j and v_j are assigned to different teams, you have to pay the price c_j . You can assume all a_i, b_i, c_j are positive.

Use a network flow algorithm to find an optimal way to divide these people into teams A and B , i.e. paying the minimum total cost. You can use any network flow algorithm, and as long as the graph you construct is polynomial-sized in n and m , you will receive full credit for time complexity.

Answer:

To solve the problem, we construct a flow network as follows:

Graph Construction

1. Nodes:

- Add a source node s and a sink node t .
- For each person i , create a node p_i .

2. Edges:

- Add an edge (s, p_i) with capacity a_i (cost of assigning i to team A).
- Add an edge (p_i, t) with capacity b_i (cost of assigning i to team B).
- For each pair of friends (u_j, v_j, c_j) , add an undirected edge (p_{u_j}, p_{v_j}) with capacity c_j (penalty for assigning u_j and v_j to different teams).

Correctness of the Reduction

Claim: We claim that the minimum s - t cut of the graph corresponds to the minimum cost of dividing people into teams A and B .

Proof. 1. Assigning person i to team A corresponds to "cutting" the edge (s, p_i) , incurring a cost of a_i . Similarly, assigning i to team B corresponds to cutting the edge (p_i, t) , incurring a cost of b_i . 2. If two friends u_j and v_j are assigned to different teams, the edge (p_{u_j}, p_{v_j}) is cut, incurring a penalty of c_j . 3. By the max-flow min-cut theorem, the value of the minimum s - t cut equals the maximum flow in the network, which minimizes the total cost. \square

Algorithm and Pseudocode We solve the problem by running a max-flow algorithm such as Edmonds-Karp. Here is the pseudocode:

Input: Costs a_i, b_i for each person i , penalties c_j for m friend pairs
Output: Minimum total cost of dividing people into teams

¹We assume the same pair of friends will only appear once; but a person may have multiple friends.

-
1. Construct a graph G as follows:
 - Add source node s and sink node t .
 - Add edges (s, p_i) with capacity a_i .
 - Add edges (p_i, t) with capacity b_i .
 - Add edges (p_u, p_v) with capacity c_j for all friend pairs (u, v) .
 2. Run Edmonds-Karp to compute the maximum flow.
 3. The minimum cost is the value of the minimum s - t cut, which is the same as the maximum flow.

Return minimum cost.

Complexity Analysis

1. Graph Size: - There are $n + 2$ nodes (one for each person, plus s and t). - There are $2n + m$ edges:

- n edges from s to p_i ,
- n edges from p_i to t ,
- m edges between pairs of friends.

Therefore, the graph size is polynomial in n and m .

2. Time Complexity: - Using Edmonds-Karp, the time complexity is $O(VE^2)$, where V is the number of vertices and E is the number of edges. - In our case, $V = O(n)$ and $E = O(n + m)$, so the complexity becomes:

$$O((n)(n + m)^2).$$

- Since the graph is polynomial-sized in n and m , this algorithm satisfies the time complexity requirements.

3. (33 points) **Meet in the Middle.** We know the knapsack problem is NP-hard when volumes are allowed to be very large. In this question, you will consider a similar problem, which is also NP-hard.

You get n real numbers (positive or negative) whose absolute values may be very large. Design an algorithm running in expected time $O(2^{n/2} \times n)$ that decides whether there is a subset of numbers that sum up to 0.

Hint. Use a hash table to store all possible subset sums for the first $\frac{n}{2}$ elements.

Answer.

To solve this question efficiently, we can split the n numbers into two half, where A consist the first $\frac{n}{2}$ elements, and B consist the last $\frac{n}{2}$ elements. The key idea is to divide the problem into two smaller subset, solve each independently, then merge the results.

First, we compute all the possible subset sum of A , and store the result in a hash table,

where the key is the subset sum and the value can be anything, for example the list of subsets producing this subset sum. For elements in B , we just compute the subset sums and store them in an array. If there are subset sums equal to 0 in either A or B , we can return that there are a subset of numbers that sum up to 0 immediately. If not, for each of the computed subset sums s in B , we check if there is a complementary sum $-s$ in A using the hash table. If there is such $-s$ exist in hash table, there is a subset of numbers that sum up to 0 in the given n real numbers, else no such subset exist.

Algorithm summary

1. Split the input array *nums* into two halves:
 - 1.1 A (first $n/2$ elements) and B (remaining $n/2$ elements).
2. Compute all subset sums for A :
 - 2.1 Calculate all possible subset sums of A .
 - 2.2 Store the subset sums in a hash table, where the key is the subset sum, the value is the list of subsets that produce this sum(or anything).
3. Compute all subset sums for B :
 - 3.1 Calculate all possible subset sums of B .
 - 3.2 Store the subset sums in an array.
4. Check for subsets summing to 0:
 - 4.1 Check if 0 exists as a subset sum in either A or B .
 - 4.2 If 0 exists, return immediately that a subset exists with a sum of 0.
5. Find complementary sums:
 - 5.1 For each subset sum s in B , check if its complement $-s$ exists in the hash table of A .
 - 5.2 If such a complement exists, return that a subset exists with a sum of 0.
6. If no subset sums to 0 in either A , B , or across A and B , Return no such subset exists.

Proof of correctness

To prove completeness, we must show that the algorithm examines all possible subsets of the given n numbers and guarantees to find the subset with a sum of 0 if such a subset exists. Let $S(A)$ be all the subset sums of A and $S(B)$ be all the subset sums of B .

Case 1: The subset is entirely in A .

The algorithm explicitly computes all possible subset sums of A . If any subset of A sums to 0, then $0 \in S(A)$, and the algorithm will detect this in step 4.1.

Case 2: The subset is entirely in B .

The algorithm explicitly computes all possible subset sums of A . If any subset of B sums to 0, then $0 \in S(B)$, and the algorithm will detect this in step 4.1.

Case 3: The subset is split between in A and B .

Any subset of the given n numbers can be expressed as the union of a subset of A and a subset of B . For such a subset, the total sum is the sum of the subset sums from A and B :

$$\text{Subset Sum} = s_A + s_B$$

where $s_A \in S(A)$ and $s_B \in S(B)$

Our algorithm iterate over all $s_B \in S(B)$ and checks whether $-s_B \in S(A)$. If such s_B exist, that means:

$$s_A + s_B = 0 \implies s_A = -s_B$$

Which ensures that all possible subset formed by combining the subset of A and B are checked.

Since all subset of the given n real numbers are either:

1. Fully in A
2. Fully in B
3. Split in A and B

Our algorithm has explicitly check all the possible subsets of the given n numbers, the algorithm is complete and correct.

Time Complexity Analysis

1. Split the input real numbers into two halves takes $O(1)$
2. Compute all subset sums for A :
 - 2.1 There are $\frac{n}{2}$ elements in A , so the number of subset in A is $2^{n/2}$, there are at most n element in a subset, so calculate all possible subset sums of A takes $O(2^{n/2} \cdot n)$.
 - 2.2 Store the subset sums in a hash table takes $O(2^{n/2})$ since there is $2^{n/2}$ subset sums in A .
3. Compute all subset sums for B :
 - 3.1 There are $\frac{n}{2}$ elements in B , so the number of subset in B is $2^{n/2}$, there are at most n element in a subset, so calculate all possible subset sums of B takes $O(2^{n/2} \cdot n)$.
 - 3.2 Store the subset sums in an array takes $O(2^{n/2})$ since there is $2^{n/2}$ subset sums in B .
4. Check if 0 exists as a subset sum in either A or B takes $O(1)$.
5. Find complementary sums:

5.1 Hash table lookup takes $O(1)$ on average, there are $2^{\frac{n}{2}}$ subsets in B , so for each subset sum s in B , check if its complement $-s$ exists in the hash table of A takes $O(2^{\frac{n}{2}})$.

So the total Time Complexity is:

$$O(2^{n/2} \cdot n) + O(2^{n/2}) + O(2^{n/2}) = O(2^{n/2} \cdot n)$$

Coding Problem Record

Yuchen:

311. Sparse Matrix Multiplication Premium Solved

Medium Topics Companies

Given two **sparse matrices** `mat1` of size `m x k` and `mat2` of size `k x n`, return the result of `mat1 x mat2`. You may assume that multiplication is always possible.

Example 1:

1	0	0
-1	0	3

 \times

7	0	0
0	0	0
0	0	1

 $=$

7	0	0
-7	0	3

1.1K 7 11 Online

Submissions

Status	Language	Runtime	Memory	Notes
Accepted	Python3	43 ms	17.1 MB	

Nov 22, 2024

686. Repeated String Match Solved

Medium Topics Companies

Given two strings `a` and `b`, return the **minimum number of times you should repeat string `a`** so that **string `b` is a substring of it**. If it is impossible for `b` to be a substring of `a` after repeating it, return `-1`.

Notice: string `"abc"` repeated 0 times is `""`, repeated 1 time is `"abc"` and repeated 2 times is `"abcbc"`.

Example 1:

Input: `a = "abcd", b = "cdababcd"`
Output: 3
Explanation: We return 3 because by repeating a three times

2.6K 29 5 Online

Submissions

Status	Language	Runtime	Memory	Notes
Accepted	Python3	15 ms	17.1 MB	

Nov 21, 2024

1392. Longest Happy Prefix Solved

Hard Topics Companies Hint

A string is called a **happy prefix** if it is a **non-empty prefix** which is also a **suffix** (excluding itself).

Given a string `s`, return the **longest happy prefix** of `s`. Return an empty string `""` if no such prefix exists.

Example 1:

Input: `s = "level"`
Output: `"l"`
Explanation: `s` contains 4 prefix excluding itself (`"l"`, `"le"`, `"lev"`, `"leve"`), and suffix (`"l"`, `"el"`, `"vel"`),

1.4K 26 6 Online

Submissions

Status	Language	Runtime	Memory	Notes
Accepted	Python3	126 ms	21.8 MB	

Nov 18, 2024

438. Find All Anagrams in a String Solved

Medium Topics Companies

Given two strings `s` and `p`, return an array of all the start indices of `p`'s **anagrams** in `s`. You may return the answer in **any order**.

Example 1:

Input: `s = "cbaebabacd", p = "abc"`
Output: `[0,6]`
Explanation:
The substring with start index = 0 is "cba", which is an anagram of "abc".
The substring with start index = 6 is "bac", which is an anagram of "abc".

12.5K 117 51 Online

Submissions

Status	Language	Runtime	Memory	Notes
Accepted	Python3	27 ms	17.2 MB	

Nov 19, 2024

Botao:

Problem List

1701. Average Waiting Time

Unlocked

Topics

Companies

Hint

There is a restaurant with a single chef. You are given an array `customers`, where `customers[i] = [arrivali, timei]`:

- `arrivali` is the arrival time of the `i`th customer. The arrival times are sorted in **non-decreasing** order.
- `timei` is the time needed to prepare the order of the `i`th customer.

When a customer arrives, he gives the chef his order, and the chef starts preparing it once he is idle. The customer waits till the chef finishes preparing his order. The chef does not prepare food for more than one customer at a time. You chef prepare food for customers in the **order they were given in the input**.

Return their **average** waiting time of all customers. Solutions within 10^{-5} from the actual answer are considered accepted.

Example 1:

Input: `customers = [[1,2],[2,5],[4,3]]`
Output: `4.00000`
Explanation:
The chef starts at time 0. He starts cooking at time 1 for the first customer. He finishes his order at time 1+2=3. Now he starts cooking at time 3 for the second customer. He finishes his order at time 3+5=8. Now he starts cooking at time 8 for the third customer. He finishes his order at time 8+3=11. The average waiting time is $(2+5+3)/3 = 4$.

Editorial **Solutions**

Test Result **Submissions**

Status **Language** **Runtime** **Memory** **Notes**

Accepted Nov 20, 2024 **Python** 24 ms 48.7 MB

Problem List

2516. Take K of Each Character From Left and Right

Unlocked

Topics

Companies

Hint

You are given a string `s` consisting of the characters `'a'`, `'b'` and `'c'` and a non-negative integer `k`. Each minute, you may take either the **leftmost** character of `s` or the **rightmost** character of `s`.

Return the **minimum** number of minutes needed for you to take **at least** `k` of each character or return `-1` if it is not possible to take `k` of each character.

Example 1:

Input: `s = "aabacaaca", k = 2`
Output: `5`
Explanation:
Take three characters from the left of `s`. You now have two 'a' characters, and one 'b' character.
Take five characters from the right of `s`. You now have four 'a' characters, two 'b' characters, and two 'c' characters.
It can be proven that 5 is the minimum number of minutes needed.

Editorial **Solutions**

Test Result **Submissions**

Status **Language** **Runtime** **Memory** **Notes**

Accepted Nov 20, 2024 **Python** 1027 ms 15.7 MB

Accepted Nov 20, 2024 **Python** 1020 ms 15.6 MB

Accepted Nov 20, 2024 **Python** 1062 ms 15.6 MB

Problem List

773. Sliding Puzzle

Unlocked

Topics

Companies

Hint

On an 2×3 board, there are five tiles labeled from 1 to 5, and an empty square represented by 0. A **move** consists of choosing 0 and a 4-directionally adjacent number and swapping it.

The state of the board is solved if and only if the board is `[[1,2,3],[4,5,0]]`.

Given the puzzle board `board`, return the least number of moves required so that the state of the board is solved. If it is impossible for the state of the board to be solved, return `-1`.

Example 1:

Input: `board = [[1,2,3],[4,5,0]]`
Output: `1`
Explanation: Swap tiles 0 and 3, it will only be one move.

Editorial **Solutions**

Test Result **Submissions**

Status **Language** **Runtime** **Memory** **Notes**

Accepted Nov 20, 2024 **Python** 12 ms 11.8 MB

Wrong Answer Nov 20, 2024 **Python** N/A N/A

Problem List

1975. Maximum Matrix Sum

Unlocked

Topics

Companies

Hint

You are given an $n \times n$ integer matrix `matrix`. You can do the following operation **any** number of times:

- Choose any two **adjacent** elements of `matrix` and **multiply** each of them by `-1`.

Two elements are considered **adjacent** if and only if they share a **border**.

Your goal is to **maximize** the summation of the matrix's elements. Return the **maximum** sum of the matrix elements using the operation mentioned above.

Example 1:

Input: `matrix = [[1,-1],[-1,1]]`
Output: `4`
Explanation: The matrix is `[[1,-1],[-1,1]]`. We can choose the two elements in the first row and multiply each by `-1`. The resulting matrix is `[[1,1],[1,1]]` and the sum of its elements is 4.

Editorial **Solutions**

Test Result **Submissions**

Status **Language** **Runtime** **Memory** **Notes**

Accepted Nov 20, 2024 **Python** 61 ms 17.6 MB

Yuheng:

Description | Accepted | Editorial | Solutions | Submissions

128. Longest Consecutive Sequence

Medium | Topics | Companies

Solved

Given an unsorted array of integers `nums`, return the length of the longest consecutive elements sequence.

You must write an algorithm that runs in $O(n)$ time.

Example 1:

Input: `nums = [100,4,200,1,3,2]`
Output: 4
Explanation: The longest consecutive elements sequence is `[1, 2, 3, 4]`. Therefore its length is 4.

Description | Accepted | Editorial | Solutions | Submissions

212. Word Search II

Hard | Topics | Companies | Hint

Solved

Given an $m \times n$ board of characters and a list of strings `words`, return all words on the board.

Each word must be constructed from letters of sequentially adjacent cells, where **adjacent cells** are horizontally or vertically neighboring. The same letter cell may not be used more than once in a word.

Example 1:

o	a	a	n
e	t	a	e

Description | Accepted | Editorial | Solutions | Submissions

148. Sort List

Medium | Topics | Companies

Given the `head` of a linked list, return the list after sorting it in **ascending order**.

Example 1:

```
graph LR; 4((4)) --> 2((2)); 2 --> 1((1)); 1 --> 3((3));
```

45. Jump Game II

Medium | Topics | Companies

You are given a **0-indexed** array of integers `nums` of length `n`. You are initially positioned at `nums[0]`.

Each element `nums[i]` represents the maximum length of a forward jump from index `i`. In other words, if you are at `nums[i]`, you can jump to `i + j` where:

- $0 \leq j \leq \text{nums}[i]$ and
- $i + j < n$

Peng Yao:

32. Longest Valid Parentheses

Given a string containing just the characters '(' and ')', return the length of the longest valid (well-formed) parentheses *substring*.

12.6K	80	13 Onl
Submissions		
Status	Language	Runtime
Accepted	C++	0 ms
a few seconds ago		10.3 MB

132. Palindrome Partitioning II

Given a string *s*, partition *s* such that every *substring* of the partition is a *palindrome*. Return the *minimum* cuts needed for a *palindrome partitioning* of *s*.

5.6K	56	6 Onl
Submissions		
Status	Language	Runtime
Accepted	C++	55 ms
a few seconds ago		13.9 MB

87. Scramble String

We can scramble a string *s* to get a string *t* using the following algorithm:

- If the length of the string is 1, stop.
- If the length of the string is > 1, do the following:
 - Split the string into two non-empty substrings at a random index, i.e., if the string is *s*, divide it to *x* and *y* where *s* = *x* + *y*.
 - Randomly decide to swap the two substrings or to keep them in the same order. i.e., after this step, *s* may become *y* + *x* or *x* + *y* or *s* = *y* + *x*.
 - Apply step 1 recursively on each of the two substrings *x* and *y*.

Given two strings *s1* and *s2* of the same length, return *true* if *s2* is a scrambled string of *s1*, otherwise, return *false*.

3.4K	82	6 Onl
Submissions		
Status	Language	Runtime
Accepted	C++	135 ms
a few seconds ago		34.7 MB

115. Distinct Subsequences

Given two strings *s* and *t*, return the number of distinct *subsequences* of *s* which equals *t*. The test cases are generated so that the answer fits on a 32-bit signed integer.

6.8K	118	27 Onl
Submissions		
Status	Language	Runtime
Accepted	C++	27 ms
a few seconds ago		13.1 MB

Bowen Yu:

